BRANDLENS

CODESTAR 30 / ROUND 2

Evidence-first brand review.

Human-led decisions, from draft to sign-off.

UPDATED PROJECT SUBMISSION | 26 OCTOBER 2026

A pre-publication checkpoint for every brand.

BrandLens helps a team check marketing copy against an editable brand policy, see exactly why a phrase needs attention, and record a human review decision. Automated checks never approve or publish a draft.

TEAM	TEAM LEADER	PRODUCT STAGE
BrandLens	Yuvraj Sonkar (solo participant)	Round 2 working build

Submission links: The app and source must be published to public URLs before this PDF is filed. This local build includes a Render deployment blueprint; no public URL is claimed until deployment and repository publication are completed.

01 / The judge's core workflow

01 Set policy	Add banned phrases and approved facts in Brand policy.
02 Submit a draft	Paste a campaign title and marketing copy; a sample is ready on first open.
03 Inspect the evidence	Review exact quoted text, rule/fact source IDs, severity, reason, and next step.
04 Make the call	Approve or request changes, add a note, and revisit the saved audit/history.

Judge tip: Run the pre-filled sample first. It intentionally triggers banned-term and unsupported-claim findings so the explanation and decision path are visible immediately.

02 / Completed Round 2 features

The implementation prioritises a complete, testable text-review loop over unbuilt screens or claims of AI-powered compliance.

Editable brand policy	Add and persist banned phrases and approved facts in SQLite.	Reviews use the team's current reference list.
Explainable draft checks	Case-insensitive exact-term checks plus numeric/absolute-claim heuristics.	Reviewers can inspect a quoted trigger and cited source ID.
Review findings	Show severity, confidence label, reason, source text, and suggested next step.	Feedback is tied to draft evidence, not an unexplained score.
Human decision	Approve or request changes with an optional reviewer note.	The system does not change or approve a draft automatically.
Review history & audit	Persist drafts, findings, status changes, notes, and dated audit events.	Reopen reviews and see who made a decision in this demo workspace.
Demo-ready interface	Sample draft, policy/history pages, mobile layout, error and empty states.	A first-time judge can try the main flow without a walkthrough.

WHAT IS NOT IN THIS BUILD

No external LLM, retrieval/vector index, document/PDF upload, image review, role-based login, tracked edits, publishing integration, or analytics dashboard. The rules are deterministic heuristics and do not provide legal advice or prove that content is compliant.

Round 2 validation evidence

Four backend workflow tests pass, covering configurable policy, quoted findings, fact matching, human decisions, saved audit events, and blank input. Frontend lint and production build pass. The running browser/API workflow was manually checked against the sample draft.

03 / Technology & implementation

Frontend	React 19, Vite, JavaScript, CSS	Draft editor, policy controls, findings, decisions, and history views.
Backend API	Python 3.11, FastAPI, Pydantic	Validated endpoints, deterministic checks, and review workflow.
Storage	SQLite	Policy, draft, result, decision, and audit event records.
Deployment	Docker multi-stage image; Render Blueprint	One web service builds the React assets and serves the API and app.
Quality checks	unittest, Oxlint, Vite build	Repeatable backend workflow checks and production frontend verification.

Evidence-first review pipeline

Draft text > **active policy** > **deterministic match** > **finding with exact quote + source ID** > **human decision + audit event**

A banned phrase is flagged only when it is found in the submitted text. Numeric/absolute claims are surfaced for reviewer verification unless an approved fact contains an exact matching claim. These safeguards are rule-based, not generative AI or a legal guarantee.

Responsible demo use

The public-demo configuration has no authentication and permits anonymous visitors. On the included free Render setup, SQLite is stored under /tmp and can reset when the instance restarts or redeploys. Do not enter confidential or personal material. Use persistent storage and add authentication before handling real organisation data.

04 / Plan changes & Round 3 improvements

Changes from the Round 1 proposal

Kept	React frontend, Python FastAPI backend, SQLite for the demo, explainable checks, human approval, and audit/history.
Focused for Round 2	Completed the end-to-end text-review loop first: starter policy, configurable rules/facts, findings, reviewer note, human decision, and saved history.
Deferred openly	Document upload, image checks/OCR, multi-user roles, external AI and semantic retrieval, tracked-change editing, and insight analytics.

Three substantive additions planned for Round 3

01	**Image brand review** Accept image assets, check brand palette and logo presence, capture required alt text, and run OCR so image text can enter the same explainable text checks.
02	**Guideline-health insights** Aggregate review findings and reviewer overrides to surface recurring problem areas and rules that may be unclear; let brand owners inspect the source before revising it.
03	**Verifier reliability upgrade** Create a seeded-error evaluation set, reject quotes absent from the submitted draft and IDs absent from policy, test false positives/regressions, and tune when uncertain claims require human verification.

Before filing: publish the source as a public repository, deploy the Docker service, open the live app link in a private window, verify that a first-time visitor can run the sample, then replace the pending links in this PDF with the final public URLs.

Project files: README.md | Dockerfile | render.yaml | backend/ | frontend/